

Supervised Learning (COMP0078) Report

Yoga Advait Veturi zchayav@ucl.ac.uk

Maki Egawa zcabmeg@ucl.ac.uk

January 17, 2021

1 PART I

The source code for this section is provided in `p1_sourcecode.ipynb`, which is attached with the submission, or is accessible via this clickable [link](#). For convenience however, relevant code snippets are also attached inline with the explanations.

1.1 Linear Regression

- (a) We are provided with a toy data set

$$\mathcal{D} = \{(1, 3), (2, 2), (3, 0), (4, 5)\}$$

where each pair (x_i, y_i) contains an input $x_i \in \mathbb{R}^n$ and output $y_i \in \mathbb{R}$. The goal is to fit \mathcal{D} using a linear combination of polynomial basis functions

$$\{\phi_1(x) = 1, \phi_2(x) = x^2, \phi_3(x) = x^3, \dots, \phi_k(x) = x^{k-1}\}$$

of dimension k , where each basis function $\phi_k : \mathbb{R}^n \rightarrow \mathbb{R}$. The fitting is performed for dimensions $k = 1, 2, 3, 4$.

To perform the curve fitting, we first defined a function Φ that takes the vector of inputs \mathbf{x} and transforms it into an $(m \times k)$ feature map $\Phi(\mathbf{x})$

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_m \end{bmatrix} \longrightarrow \Phi(\mathbf{x}) = \begin{bmatrix} \phi_1(x_1) & \dots & \phi_k(x_1) \\ \phi_1(x_2) & \dots & \phi_k(x_2) \\ \vdots & \ddots & \vdots \\ \phi_1(x_m) & \dots & \phi_k(x_m) \end{bmatrix}$$

where each $\phi_k(x)$ is a single polynomial basis function. The following code executes this, provided the value of k .

```

1 #Computes the polynomial basis
2 def phi(x, k):
3     """ Computes a single basis function  $x^{k-1}$ , given an input
4         vector x of shape (m, 1) and the value k. m is the number of
5         training examples """
6     return x**(k-1)
7
8 #This function creates the feature map PHI()
9 def PHI(x, k):
10     """ Creates the feature map, which is a matrix of shape (m, k).
11         Each column is the kth basis function applied on
12         the input vector x of shape (m, 1). m is the number of training
13         examples """
14
15     #Number of training examples
16     m = x.shape[0]
17
18     #Initialize a feature map of shape (m, k) with 0s in it.
19     feature_map = np.zeros((m, k))
20
21     #In every column, replace the zeros with the basis function k
22     #computed on input vector x
23     for i in range(0, k):
24         feature_map[:, i] = phi(x.squeeze(), i+1)
25
26     return feature_map
27
28 #Compute feature maps for k = {1,...,4}
29 x_tr_k1 = PHI(x, k=1)
30 x_tr_k2 = PHI(x, k=2)
31 x_tr_k3 = PHI(x, k=3)
32 x_tr_k4 = PHI(x, k=4)

```

With the transformed feature map $\Phi(\mathbf{x})$, we then computed the weights $\mathbf{w} \in \mathbb{R}^k$ that minimises the sum of squared errors defined by

$$SSE = \sum_{i=1}^m (y_i - \mathbf{w} \cdot \phi(x))^2 = (\Phi \mathbf{w} - y)^T (\Phi \mathbf{w} - y)$$

Setting the gradient $\nabla_{\mathbf{w}}(SSE) = 0$ and solving for \mathbf{w} , the least squares solution is

$$\mathbf{w} = (\Phi^T \Phi)^{-1} \Phi^T y$$

The following code computes the weights for $k = 1, 2, 3, 4$.

```
1 #Compute weights for the transformed data using least squares
  regression
2 w_k1 = np.linalg.inv(x_tr_k1.T @ x_tr_k1) @ x_tr_k1.T @ y
3 w_k2 = np.linalg.inv(x_tr_k2.T @ x_tr_k2) @ x_tr_k2.T @ y
4 w_k3 = np.linalg.inv(x_tr_k3.T @ x_tr_k3) @ x_tr_k3.T @ y
5 w_k4 = np.linalg.inv(x_tr_k4.T @ x_tr_k4) @ x_tr_k4.T @ y
```

(b) The following polynomial equations were obtained, corresponding to $k = 1, 2, 3, 4$:

$$y = 2.5$$

$$y = 1.5 + 0.4x$$

$$y = 9.0 - 7.1x + 1.5x^2$$

$$y = -5.0 + 15.17x - 8.5x^2 + 1.33x^3$$

These are visualized in Figure 1.

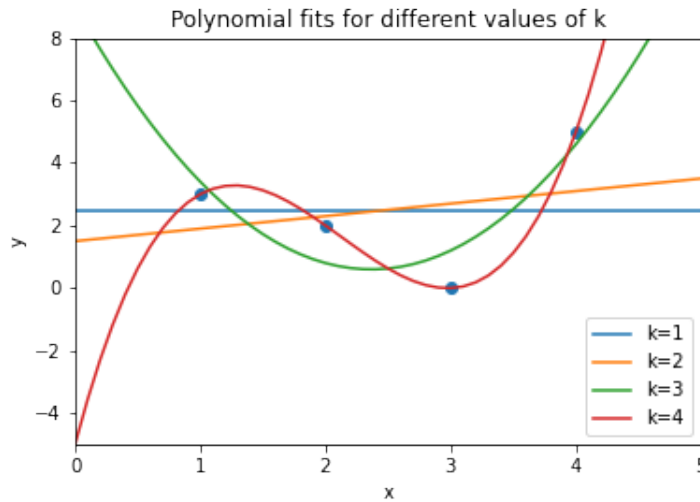


Figure 1: Fitted curves for $k=1,2,3,4$

(c) For each fitted curve, the mean squared error was computed using $MSE = \frac{SSE}{m}$. The following code executes the MSE function.

```
1 def MSE(x_values, y_values, weights):
2     """
3     Computes the MSE given the x_values matrix (m, n), y_values
4     vector of shape (m, 1) and weights vector of shape (n, 1).
5     m is the number of training examples, n is the number of input
6     features.
7     """
8     m = x_values.shape[0]
9     SSE = ((x_values @ weights) - y_values).T @ ((x_values @
10     weights) - y_values)
11     MSE = SSE / m
12     return MSE
```

```

10
11 #Print MSE values for each of the polynomial bases
12 print("MSE for k = 1 : " + str(MSE(x_tr_k1, y, w_k1).item()))
13 print("MSE for k = 2 : " + str(MSE(x_tr_k2, y, w_k2).item()))
14 print("MSE for k = 3 : " + str(MSE(x_tr_k3, y, w_k3).item()))
15 print("MSE for k = 4 : " + str(MSE(x_tr_k4, y, w_k4).item()))

```

The MSE values are presented in Table 1 below.

k	MSE
1	3.25
2	3.05
3	0.80
4	3.49e-23

Table 1: Mean Squared Error for fitted curves $k = 1, 2, 3, 4$

2. We now fit models on data generated from the random function

$$g_{\sigma=0.07}(x) := \sin^2 2\pi x + \epsilon$$

where ϵ is Gaussian noise with a mean of 0 and variance $\sigma^2 = 0.07^2$.

- (a) i. A training set of 30 points was generated by sampling random points uniformly in the interval $[0, 1]$. These points were passed through g_σ and on each call of the function, Gaussian noise ϵ is added. The following code implements this:

```
1 def g_sigma(x, sigma):
2     """
3     Random function of form g_sigma(x) = sin^2(2*pi*x) + epsilon
4     where epsilon is a Gaussian distributed random variable
5     with mean 0 and sigma^2 variance.
6
7     Takes an input x and returns the value g_sigma(x) with an
8     added noise value
9     generated from epsilon.
10    """
11
12    #Compute function on the point x
13    function = np.square(np.sin(2 * np.pi * x))
14    #Generate a single noise value for that point x
15    noise = np.random.normal(0.0, sigma)
16
17    return function + noise
18
19 def generate_data(interval, size, sigma):
20     """
21     Generates a random sample {(x_1, g_1), ..., (x_m, g_m)}. First
22     samples a single x value from the interval
23     and then computes g_sigma = sin^2(2pi*x) + epsilon, where
24     epsilon is random gaussian noise.
25
26     Args
27     ----
28     interval - tuple of (lower_bound, upper_bound)
29     size - number of points to generate
30     sigma - standard deviation of the noise
31
32     Returns
33     -----
34     NumPy arrays of the generated input and output data
35     """
36
37     #This will store the x_values and g_sigma(x) values
38     x_values = []
39     g_values = []
40
41     for i in range(size):
42
43         #Generate a single point
44         x = np.random.uniform(interval[0], interval[1])
```

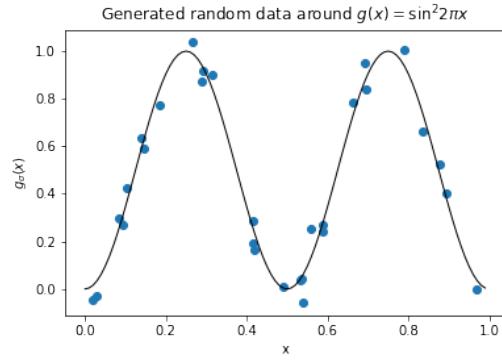
```

41     #Create the output value using the random function
42     g = g_sigma(x, sigma=sigma)
43     #Append the values to the list
44     x_values.append(x)
45     g_values.append(g)
46
47     return np.array(x_values), np.array(g_values)
48
49 #Generate training data - set a seed for ensuring reproducible
    results
50 np.random.seed(1)
51 x_train, g_train = generate_data(interval=(0, 1), size=30,
    sigma=0.07)

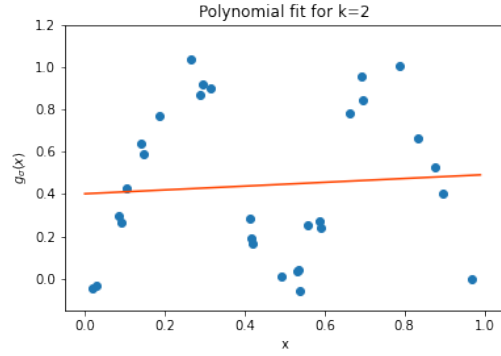
```

The generated data set is plotted along with the deterministic function $\sin^2(2\pi x)$ in Figure 2a.


- ii. This data set is then fit using polynomial bases of dimension $k = 2, 5, 10, 14, 18$ using the same process followed in 1a. The x values are first transformed into the feature map Φ , then using least squares regression, the weights \mathbf{w} are computed for each k . These fitted curves are presented in Figure 2(b-f).
- (b) The training error denoted te_k was computed by performing linear regression on the training points for every value of $k = 1, \dots, 18$ and calculating the MSE on the training points using the $MSE()$ function. The natural log of the training error is plotted in Figure 3(a). It is noticeable that the function is decreasing, however around epoch 16, there is a rise, which could be attributed to numerical errors.
- (c) A test set of 1000 points was generated using the *generate_data()* function described earlier. The test set error, denoted tse_k was computed by calculating the MSE on the test points using the weights computed on the training points. The natural log of the test set error for $k = 1, \dots, 18$ is plotted in Figure 3(b).
- (d) The previous results correspond to one type of generated training and test set. To see these errors smoothed out for different types of training and test sets, the previous steps in 2b and 2c were repeated for 100 runs and the log of the average train and test set error over the runs were computed. The results are plotted in Figure 3(c).




(a)



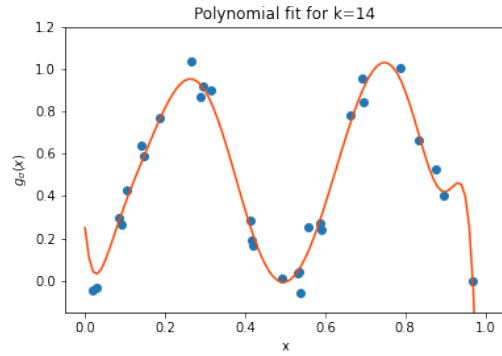
(b)




(c)



(d)



(e)



(f)

Figure 2: (a) Generated data (blue points) with the deterministic function $\sin^2(2\pi x)$ overlapped (black curve). All fitted curves are shown in orange. (b) Fitted curve for $k = 2$. (c) Fitted curve for $k = 5$. (d) Fitted curve for $k = 10$. (e) Fitted curve for $k = 14$. (f) Fitted curve for $k = 18$.



Figure 3: MSE values for $k = 1, \dots, 18$ (a) Log training error (b) Log test set error (c) Log training and test set error, averaged over 100 runs

3. The experiments performed in 2b, 2c and 2d are repeated, now using a different basis

$$\{\sin 1\pi x, \sin 2\pi x, \dots, \sin k\pi x\}$$


We modify the functions for computing the feature map as follows:

```

1 #Computes the basis function
2 def phi_v2(x, k):
3     """ Computes a single basis function, given input vector x of shape
4         (m, 1) and value k. m is the number of training examples """
5     return np.sin(k*np.pi*x)
6
7 #This function creates the feature map
8 def PHI_v2(x, k):
9     """ Creates the feature map, which is a matrix of shape (m, k).
10        Each column is the kth basis function applied on
11        the input vector x of shape (m, 1). m is the number of training
12        examples """
13
14    #Extract the number of training examples
15    m = x.shape[0]
16
17    #Initialize feature map matrix to matrix of zeros
18    feature_map = np.zeros((m, k))
19
20    #Iterative over each column
21    for i in range(0, k):
22        #For each column of the feature map matrix, replace the 0s with
23        #the computed kth basis function values
24        feature_map[:, i] = phi_v2(x.squeeze(), i+1)
25    return feature_map

```


The log te_k and tse_k for $k = 1, \dots, 18$ are plotted in Figure 4a and 4b respectively. The entire operation is performed for 100 runs, using different generated training and test points. The log of the average te_k and tse_k is visualized in Figure 4c.



(a)



(b)



(c)

Figure 4: (a) Log training error for $k = 1, \dots, 18$ (b) Log test error for $k = 1, \dots, 18$ (c) Log of average train and test error over 100 runs.

1.2 Filtered Boston Housing and Kernels

4. In this section, we were provided with a modified version of the Boston Housing data set. Below is a quick inspection of the first few rows of this data.

CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	LSTAT	MEDV
0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	4.98	24.0
0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	9.14	21.6
0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	4.03	34.7
0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	2.94	33.4
0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	5.33	36.2

Table 2: Inspection of first 5 rows of Boston Housing data set

The goal was to predict the median house price (MEDV) using the remaining 12 variables. To do so, we performed multiple experiments with different variants of linear regression. These experiments involved splitting the data into a training set, which contained 2/3 of the examples, and a test set with 1/3 of the examples. The following function implements the data splitting. Note that the data can be shuffled before splitting in order to generate random splits.

```
1 def split_data(inputs, targets, test_proportion, shuffle=None):
2     """
3     Splits the data into training and test sets.
4
5     Args
6     ----
7     inputs : NumPy array of input data. Should be of shape (# examples,
8                 # features)
9     targets : NumPy array of target data. Should be of shape (#
10                examples, 1)
11     test_proportion : Value between 0 and 1 which specifies how much of
12                        the data to use for testing.
13     shuffle : Optional. Set to True if you want the data shuffled and
14                then split.
15     seed : Optional. Set for reproducible results.
16
17     Returns
18     -----
19     train_X : NumPy array of training examples. Should be of shape (#
20                examples, # features)
21     train_Y : NumPy array of training targets. Should be of shape (#
22                examples, 1)
23     test_X : NumPy array of testing examples. Should be of shape (#
24                examples, # features)
25     test_Y : NumPy array of testing targets. Should be of shape (#
26                examples, 1)
27     """
28
29     #Stores the number of data points
30     nData = inputs.shape[0]
31
32     # Shuffle data
```

```

25  if shuffle:
26      #Generate a shuffled version of the array indices
27      shuffled_indices = np.random.permutation(nData)
28      #Shuffle the inputs as per in the array of shuffled indices
29      shuffled_inputs = inputs[shuffled_indices, :]
30      shuffled_targets = targets[shuffled_indices, :]
31  else:
32      #If shuffle is set to False then we just work with the data in
      its original order
33      shuffled_inputs = inputs
34      shuffled_targets = targets
35
36  # Calculate the split index based on the specified proportions
37  split_index = int((1 - test_proportion) * nData)
38
39  # Select the examples up to the split index to be used as training
      set
40  train_X = shuffled_inputs[:split_index]
41  train_Y = shuffled_targets[:split_index]
42  # Select the examples from the split index onwards to be used as
      the test set
43  test_X = shuffled_inputs[split_index:]
44  test_Y = shuffled_targets[split_index:]
45
46  return train_X, train_Y, test_X, test_Y

```

- (a) We first performed Naive regression. The data set was split into train and test sets using the splitting function. A vector of ones was created with the same length as the training set and another for the test set. Using the training ones vector and the training outputs, we computed the weights with linear regression. Finally, the train and test MSE was computed using the $MSE()$ function. The code implementation for this is provided below. Note that it is performed over 20 runs for 20 random train/test splits.

```
1 #For reproducibility of results
2 np.random.seed(101010)
3
4 #Stores the number of runs
5 runs = 20
6
7 #These will store the total train and test MSE over the 20 runs
8 train_MSE = 0
9 test_MSE = 0
10
11 #Iterate over runs
12 for i in range(runs):
13
14     #Generate samples
15     xtrain, ytrain, xtest, ytest = split_data(inputs=X, targets=Y,
16         test_proportion=1/3, shuffle=True)
17
18     train_ones = np.ones((xtrain.shape[0], 1))
19     test_ones = np.ones((xtest.shape[0], 1))
20
21     #Perform Naive Regression
22     w_naive = np.linalg.inv(train_ones.T @ train_ones) @ train_ones
23         .T @ ytrain
24
25     #Calculate and aggregate train and test MSE
26     train_MSE += MSE(train_ones, ytrain, w_naive)
27     test_MSE += MSE(test_ones, ytest, w_naive)
28
29     print(">> Run {} Completed.".format(i+1))
```

The average train and test MSE over 20 runs is as follows:

Average Train MSE = 86.05624289198637

Average Test MSE = 81.51650453132758

- (b) With Naive Regression, we are essentially fitting the data using a constant function and asking the question "What would the model predict in the absence of any features?". The interpretation of this constant function is as follows:

We can define the training inputs and outputs as the vectors X_{train} and Y_{train}

$$X_{train} = \begin{pmatrix} 1 \\ 1 \\ \vdots \\ 1 \end{pmatrix}, Y_{train} = \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{pmatrix}$$

where m is the number of training examples. If we used these vectors to compute the weights w_{naive} then we would get the following solution:

$$w_{naive} = (X_{train}^T X_{train})^{-1} X_{train}^T Y_{train} = \frac{1}{m} \sum_{i=1}^m y_i$$

Hence, the naive weights is essentially the mean of the training output values, which is what the model would predict for a given input. To validate this, we compared the output of the w_{naive} with the mean of Y_{train} on a single run and found them to be the same:

w_naive value = [[23.11780415]]

Mean of Y_train values = 23.11780415430267

- (c) We next performed Linear Regression over single attributes. A column of ones was added to the matrix of data in order to account for the bias term in the linear regression. The following function adds the ones column:

```
1 #Function to add ones to the training set as a way to account for
  the bias term in the weights
2 def add_bias(data):
3     """
4     Adds a column of ones into the dataset in order to account for
      the bias term.
5
6     Args
7     ----
8     data : a NumPy array of the input data. Can be a matrix of
      shape (m, n) or a single input vector x_i of shape (m, 1). m
      is the number of training examples
9     while n is the number of input features.
10
11     Returns
```

```

12 -----
13 data_with_bias : a NumPy array of data with 1s added. If data
14 is a matrix, then a column of 1s is added to the first column.
15 If data is a single vector x_i, then the first element of x_i
16 is 1.
17 """
18
19 #If the data is a vector
20 if data.shape[1] == 1:
21     #Create a one
22     ones = np.ones((data.shape[1], ))
23     #Insert 1 as the first value for the vector
24     data_with_bias = np.insert(data, 0, ones, axis=0)
25
26 #If data is a matrix
27 else:
28     #Create a vector ones
29     ones = np.ones((data.shape[0], ))
30     #Stack this vector of ones to the beginning of the matrix
31     data_with_bias = np.insert(data, 0, ones, axis=1)
32
33 return data_with_bias
34
35 #Convert the pandas dataframe into matrix
36 data = dataframe.values
37
38 #Add a column of 1's to the data matrix to account for bias term
39 data_with_bias = add_bias(data)
40
41 #Split the data matrix into a matrix of the inputs and the vector
42   of outputs
43 X = data_with_bias[:, :13]
44 Y = data_with_bias[:, 13].reshape(-1, 1)

```

Using the *split_data()* function, the data was shuffled and then split into training and testing sets. For training, the individual columns of the training data matrix, which represent the different attributes, were selected along with the "ones" column and the weights were computed using linear regression for these "sub-matrices". The training and test MSE was then computed using the *MSE()* function. We report the results averaged over 20 runs for 20 random train/test splits in Table 3 below.

- (d) The next experiment was linear regression over all attributes together. We followed the exact same procedure as in the single attributes, except now we computed the weights using the entire training data matrix. We report the training and test MSE averaged over 20 runs for 20 random train/test splits as follows.

Average Train MSE = 21.414764727427052
Average Test MSE = 25.922665880996654

It can be seen that our train and test MSE for linear regression over all attributes is smaller than the regressors for single attributes. This is expected as our models

Method	MSE Train	MSE Test
Attribute 1	69.568609	77.953024
Attribute 2	72.015690	76.825691
Attribute 3	62.422693	69.677222
Attribute 4	80.442598	85.280471
Attribute 5	67.203539	73.004103
Attribute 6	42.845189	45.544419
Attribute 7	70.498032	76.720579
Attribute 8	77.077643	83.831789
Attribute 9	70.034179	76.850144
Attribute 10	63.711957	70.728458
Attribute 11	61.489686	65.364829
Attribute 12	37.487598	40.817604

Table 3: Train and Test MSE for Linear regression over single attributes. Averaged over 20 runs.

incorporated more features, allowing for better prediction of the output values.

1.3 Kernelized Ridge Regression

5. In this section, we perform kernel ridge regression over the training set with ℓ examples. This involved computing the dual solution $\boldsymbol{\alpha}^*$ that optimises the equation

$$\boldsymbol{\alpha}^* = \arg \min_{\boldsymbol{\alpha} \in \mathbb{R}^\ell} \frac{1}{\ell} \sum_{i=1}^{\ell} \left(\sum_{j=1}^{\ell} \alpha_j K_{i,j} - y_i \right)^2 + \gamma \boldsymbol{\alpha}^T \mathbf{K} \boldsymbol{\alpha} \quad (1)$$

where $K(x_i, x_j)$ (abbreviated $K_{i,j}$) is the Gaussian kernel computed using x_i and x_j :

$$K(x_i, x_j) = \exp \left(-\frac{\|x_i - x_j\|^2}{2\sigma^2} \right) \quad (2)$$

\mathbf{K} is an $(\ell \times \ell)$ matrix containing the entries $K_{i,j}$.

The dual solution is thus defined as

$$\boldsymbol{\alpha}^* = (\mathbf{K} + \gamma \ell \mathbf{I}_\ell)^{-1} \mathbf{y} \quad (3)$$

where γ is a regularization parameter, \mathbf{I}_ℓ is an $(\ell \times \ell)$ identity matrix and \mathbf{y} is the vector of training outputs.

To evaluate the regression function on a test point, the formula is

$$y_{test} = \sum_{i=1}^{\ell} \alpha_i^* K(x_i, x_{test}) \quad (4)$$

- (a) In order to perform kernel ridge regression, the data set was split into training and testing sets (2/3:1/3) using the splitting function. As can be seen in equations 1 and 2, there are two additional hyperparameters σ and γ . In order to select the best values for these hyperparameters, we first performed kernel ridge regression on the training set using five-fold cross-validation. The chosen range of values for σ and γ were $[2^7, 2^{7.5}, \dots, 2^{13}]$ and $[2^{-40}, 2^{-39}, \dots, 2^{-26}]$ respectively.

The following are the steps for 5-fold cross-validation on a single pair of (σ, γ) . We performed this over all pairs of hyperparameters.

- i. Shuffle the training set. (To reduce bias)
- ii. Divide into $K=5$ groups.
- iii. Select first group as the "validation" set and use the remaining $K-1$ groups as the training set.
- iv. Compute the dual solution $\boldsymbol{\alpha}^*$ using the training set.
- v. Compute predictions for validation data.
- vi. Compute MSE for validation set and record value.
- vii. Repeat steps (iii-vi) on all groups.

In order to perform steps (iv) and (v), three components were crucial: the computation of the kernel matrix \mathbf{K} , the dual solution $\boldsymbol{\alpha}^*$ and the evaluation function (4). The code for these three components are presented below. Note that with the computation of \mathbf{K} and the evaluation function, the simple method would be to use for-loops. However, this can be computationally expensive, hence we developed a vectorized implementation, for which a detailed explanation is provided in Appendix 2.1.

```

1 def gaussian_kernel_matrix(X, sigma):
2     """
3     Computes the gaussian kernel matrix on a given set of vectors.
4
5     Args
6     ----
7     X : NumPy array of data. Should be of shape (l, n) where l is
8         number of examples and n is number of features
9     sigma : sigma value of Gaussian kernel.
10
11     Returns
12     -----
13     K : a symmetric positive definite matrix of shape (l, l)
14     """
15     #Number of samples
16     l = X.shape[0]
17
18     # Extract the diagonals of  $XX^T$  into a vector and stack it l
19     # times into lxl matrix B
20     inner_product_vector = np.diagonal(X@X.T).reshape(-1, 1)
21
22     #Compute matrix of just the dot products between  $x_i$  and  $x_j$ 
23     inner_product_matrix = inner_product_vector - 2 * (X@X.T) +
24         inner_product_vector.T
25
26     #Compute gaussian kernel matrix
27     K = np.exp((-1/(2*sigma**2)) * inner_product_matrix)
28
29     #Checkpoint to ensure matrix K is correct dimensions
30     assert K.shape == (l, l), "K matrix should be of shape ({}, {})"
31         ".format(l, l)
32
33     return K
34
35 def kernel_ridge_regression(X, y, sigma, gamma):
36     """
37     Compute the solution  $\alpha^*$  for kernelized ridge regression.
38
39     Args
40     ----
41     X : NumPy array of data. Should be of shape (l, n) where l is
42         number of examples and n is number of input features
43     y : NumPy array of target outputs. Should be of shape (l, 1)
44     sigma : sigma value for Gaussian kernel
45     gamma : regularization parameter for ridge regression
46
47     Returns
48     -----
49     alpha_star : (l, 1) vector which represents the dual solution
50         for ridge regression.
51     """
52
53     #Extract number of samples

```

```

50 l = X.shape[0]
51
52 #Compute kernel
53 K = gaussian_kernel_matrix(X, sigma) #kernel_matrix(X, X, sigma
    =sigma)
54
55 #Create identity matrix used in the formula
56 I = np.eye(l)
57
58 #Compute alpha_star value using formula
59 try:
60     alpha_star = np.linalg.inv(K + gamma*l*I) @ y
61 except np.linalg.LinAlgError:
62     alpha_star = np.linalg.pinv(K + gamma*l*I) @ y
63
64 #Checkpoint to ensure the alpha is correct dimensions
65 assert alpha_star.shape == (l, 1), "Alpha star should be of
    shape ({}, {})".format(l, 1)
66
67 return alpha_star
68
69 def evaluate(alpha, X_tr, X_te, sigma):
70     """
71     Vectorized implementation of kernel ridge regression evaluation
        function on a test point.
72
73     Args
74     ----
75     alpha : NumPy array of the dual solution alpha. Should be of
        shape (l, 1) where l is the number of training examples.
76     X_tr : NumPy array of the training set inputs. Should be of
        shape (l, n)
77     X_te : NumPy array of the test set inputs. Should be of shape (
        m, n) where m is number of test examples.
78     sigma : Sigma value of the Gaussian kernel.
79
80     Returns
81     -----
82     pred : NumPy array of predictions for the test data. Should be
        of shape (m, 1)
83     """
84
85     #Extract dimensions of data
86     l = X_tr.shape[0]
87     m = X_te.shape[0]
88
89     #Evaluate regression model at all test points
90     inner_product_matrix = np.diagonal(X_tr@X_tr.T).reshape(1, -1)
        - 2 * (X_te @ X_tr.T) + np.diagonal(X_te@X_te.T).reshape(-1,
        1)
91     K = np.exp((-1/(2*sigma**2)) * inner_product_matrix)
92     pred = K @ alpha
93
94     assert pred.shape == (m, 1), "Predictions should be of shape
        ({}, 1)".format(m, 1)

```

```

95
96     return pred

```

The final implementation of the entire 5-fold cross-validation is as follows:

```

1 def train_with_kfoldCV(k, data, kernel_params, shuffle=True,
2     verbose=True):
3     """
4     Performs k-fold cross validation on a given dataset.
5
6     Args
7     ----
8     k - number of folds of cross-validation
9     data - tuple of training input and output data. Training input
10         must be NumPy array of shape (# examples, # features) while
11         output must be of shape (# examples, 1)
12     kernel_params - tuple of (sigma, gamma) parameters for
13         computing the Gaussian kernel.
14     shuffle - If True, then data is shuffled
15     seed - to ensure reproducible results
16     verbose - Set to True if you wish to see printed notifications.
17
18     Returns
19     -----
20     cv_error - the validation set error averaged over the k folds
21         of cross validation for the specified kernel parameters
22     """
23
24     #Extract data
25     x, y = data
26
27     #Extract gaussian kernel params
28     sigma, gamma = kernel_params
29
30     #Extract dimensions
31     nSamples, nFeatures = x.shape
32
33     #Shuffle dataset randomly
34     if shuffle:
35         perm = np.random.permutation(nSamples)
36         x_shuffled = x[perm, :]
37         y_shuffled = y[perm, :]
38     else:
39         x_shuffled = x
40         y_shuffled = y
41
42     # print("X_shuffled = ", x_shuffled)
43
44     #Split data into k-groups
45     x_groups = np.array_split(x_shuffled, k)
46     y_groups = np.array_split(y_shuffled, k)
47
48     #Stores the cross validation MSE which is the MSE calculated
49     #over the validation sets generated during each fold of k-fold
50     #CV
51     cv_mse = 0

```

```

45
46 #Iterate over each cross-validation group
47 for i in range(len(x_groups)):
48
49     #Use the selected group as "validation" set
50     val_inputs, val_outputs = x_groups[i], y_groups[i]
51
52     #Use rest of groups as training set
53     train_inputs = np.vstack([x_groups[j] for j in range(len(
54 x_groups)) if j != i])
54     train_outputs = np.vstack([y_groups[j] for j in range(len(
55 x_groups)) if j != i])
56
57     #Fit model on train set
58     alpha_star = kernel_ridge_regression(X=train_inputs, y=
59 train_outputs, sigma=sigma, gamma=gamma)
60
61     #Generate predictions on "validation" set
62     val_predictions = evaluate(alpha=alpha_star, X_tr=
63 train_inputs, X_te=val_inputs, sigma=sigma)
64
65     #Compute validation errors
66     cv_mse += (1/len(val_outputs)) * (val_outputs -
67 val_predictions).T @ (val_outputs - val_predictions)
68     if verbose:
69         print(">> Cross-validation Fold {} Completed.".format(i+1))
70
71     #Average the errors
72     cv_error = cv_mse.item() / k
73
74     if verbose:
75         print(">> MEAN CROSS-VALIDATION ERROR = {}\n".format(cv_error
76 ))
77
78     #Summarize performance using sample of model eval scores
79     return cv_error
80
81 #Create parameter space
82 sigma_values = 2*np.arange(7, 13.5, 0.5)
83 gamma_values = (2*np.arange(40.0, 25.0, -1))**-1
84
85 #Represent all (sigma, gamma) pairs using meshgrid. This will be
86 helpful when creating the 3D plot
87 gamma_array, sigma_array = np.meshgrid(gamma_values, sigma_values
88 )
89
90 #This grid will store the cross validation error values
91 errors = np.zeros((sigma_values.shape[0], gamma_values.shape[0]))
92
93 #Split data into train and test set - for reproducibility, I've
94 set a seed
95 np.random.seed(100000)
96 X_train, Y_train, X_test, Y_test = split_data(inputs=X[:, 1:],
97 targets=Y, test_proportion=1/3, shuffle=True)
98 train_data = (X_train, Y_train)

```

```

90
91 #Train the model for all the different values of sigma and gamma
92 for i in range(errors.shape[0]):
93     for j in range(errors.shape[1]):
94
95         #Create the sigma and gamma pair
96         sigma = sigma_values[i]
97         gamma = gamma_values[j]
98
99         print("Sigma = {}, Gamma = {}".format(sigma, gamma))
100
101         #Perform 5-fold cross validation. Note - Because we are
102         #adding a shuffle operation inside crossval function, every
103         #pair of sigma and gamma may train on different examples
104         #To ensure that each (sigma, gamma) pair is used on the same
105         #set of training and test examples on each fold, I've seed a
106         #seed
107         np.random.seed(130399)
108         cv_error = train_with_kfoldCV(k=5, data=train_data,
109         kernel_params=(sigma, gamma), shuffle=True, verbose=True)
110
111         #Add this error value into our errors matrix
112         errors[i, j] = cv_error
113
114 print("Ridge Regression Completed!")

```




Figure 5: Surface plot of cross validation error as a function of σ and γ values. For convenience, the σ and γ values are plotted on a log-2 scale.

- (b) The cross-validation error as a function of the σ and γ values is plotted as a surface plot in Figure 5.
- (c) The best (σ, γ) pair was defined to be the one that had the least cross validation MSE. This was found to be $\sigma = 2^{9.5}$ and $\gamma = 2^{-36.0}$. Using these values, we again performed kernel ridge regression, this time using the training set and testing set. The train and test MSE for these hyperparameters is as follows:

$$\begin{aligned}\text{Train MSE} &= 5.035038540967276 \\ \text{Test MSE} &= 16.479302486525636\end{aligned}$$

- (d) Finally, all the different experiments from 4(a,c,d) and 5(a,c) were repeated for 20 random train/test splits to smooth out the results. For convenience, all the experiments for Naive Regression, Linear Regression and Kernel Ridge Regression were converted into functions, which are presented in Appendix 2.2 of the report.

The summarized results are presented in the following table.

Method	MSE Train	MSE Test
Naive Regression	83.517 ± 3.055	86.421 ± 6.143
attribute 1	69.581 ± 2.623	77.042 ± 5.180
attribute 2	72.068 ± 3.240	76.581 ± 6.554
attribute 3	63.221 ± 3.033	67.842 ± 6.142
attribute 4	81.172 ± 3.244	84.269 ± 6.907
attribute 5	67.346 ± 3.129	72.556 ± 6.382
attribute 6	44.215 ± 2.740	42.725 ± 5.437
attribute 7	70.518 ± 3.365	76.685 ± 6.933
attribute 8	77.569 ± 3.298	82.644 ± 6.792
attribute 9	70.002 ± 3.143	76.719 ± 6.460
attribute 10	63.563 ± 3.278	70.825 ± 6.664
attribute 11	61.310 ± 2.796	65.757 ± 5.560
attribute 12	37.638 ± 1.791	40.392 ± 3.651
all attributes	21.736 ± 1.472	25.097 ± 2.926
Kernel Ridge Regression	7.229 ± 1.206	13.738 ± 1.577

Table 4: Summarized results for all regression experiments averaged over 20 runs.
Results reported as mean \pm standard deviation.

2 Appendices

2.1 Vectorized Implementation of Kernel Ridge Regression

Here, we present the mathematical formulation for the vectorized implementation of kernel ridge regression. This allows us to avoid using for-loops, thus improving efficiency of training.

2.1.1 Kernel Matrix

We first vectorize the computation of the kernel matrix \mathbf{K} , which is required to compute the dual solution α^* . To do this, we begin with the original description of \mathbf{K} :

$$\mathbf{K} = \begin{bmatrix} K(x_1, x_1) & \cdots & K(x_1, x_\ell) \\ \vdots & \ddots & \vdots \\ K(x_\ell, x_1) & \cdots & K(x_\ell, x_\ell) \end{bmatrix}$$

$$K(x_i, x_j) = \exp\left(-\frac{\|x_i - x_j\|^2}{2\sigma^2}\right) = \exp\left(-\frac{\langle x_i - x_j, x_i - x_j \rangle}{2\sigma^2}\right)$$

Note that ℓ is the number of training examples, thus \mathbf{K} is an $(\ell \times \ell)$ matrix. Using the NumPy framework, we can reformulate \mathbf{K} as:

$$\mathbf{K} = \text{np.exp}\left(-\frac{1}{2\sigma^2} \begin{bmatrix} \langle x_1 - x_1, x_1 - x_1 \rangle & \cdots & \langle x_1 - x_\ell, x_1 - x_\ell \rangle \\ \vdots & \ddots & \vdots \\ \langle x_\ell - x_1, x_\ell - x_1 \rangle & \cdots & \langle x_\ell - x_\ell, x_\ell - x_\ell \rangle \end{bmatrix}\right)$$

The goal now is to breakdown the matrix in the above expression into a combination of simpler matrices. Note that a general inner product in this matrix can be expanded as follows:

$$\langle x_i - x_j, x_i - x_j \rangle = \langle x_i, x_i \rangle - 2\langle x_i, x_j \rangle + \langle x_j, x_j \rangle \quad (5)$$

We can use this property to redefine the matrix as follows:

$$\begin{bmatrix} \langle x_1 - x_1, x_1 - x_1 \rangle & \cdots & \langle x_1 - x_\ell, x_1 - x_\ell \rangle \\ \vdots & \ddots & \vdots \\ \langle x_\ell - x_1, x_\ell - x_1 \rangle & \cdots & \langle x_\ell - x_\ell, x_\ell - x_\ell \rangle \end{bmatrix} = \mathbf{B} - 2\mathbf{X}\mathbf{X}^T + \mathbf{B}^T \quad (6)$$

where

$$\mathbf{B} = \begin{bmatrix} \langle x_1, x_1 \rangle & \cdots & \langle x_1, x_\ell \rangle \\ \vdots & \ddots & \vdots \\ \langle x_\ell, x_1 \rangle & \cdots & \langle x_\ell, x_\ell \rangle \end{bmatrix}, \mathbf{X} = \begin{bmatrix} \mathbf{x}_1^T \\ \vdots \\ \mathbf{x}_\ell^T \end{bmatrix} \quad (7)$$

Essentially what we have done is break down each element of the LHS of 6 into the 3 terms as per equation 5 and then the first, second and third terms are separated into individual matrices. The second matrix consists of the terms $-2\langle x_i - x_j, x_i - x_j \rangle$, which can be simplified to the matrix $-2\mathbf{X}\mathbf{X}^T$ where \mathbf{X} is just the $(\ell \times n)$ training data matrix.

It can also be noticed that the first and third matrices are just transposes (hence \mathbf{B} and \mathbf{B}^T) of each other and contain the diagonal terms of the $\mathbf{X}\mathbf{X}^T$ matrix stacked ℓ times as rows or columns. When implementing in NumPy however, we could just represent \mathbf{B} as a single column of the diagonals of $\mathbf{X}\mathbf{X}^T$. Due to broadcasting, we can expect that the vector and matrix additions will still work. Thus, we summarize the \mathbf{K} reformulation as

$$\mathbf{K} = \text{np.exp} \left(-\frac{1}{2\sigma^2} (\mathbf{B} - 2\mathbf{X}\mathbf{X}^T + \mathbf{B}^T) \right)$$

2.1.2 Evaluation function

The next component is vectorizing the evaluation function. This function has the form

$$y_{test} = \sum_{i=1}^l \alpha_i^* K(x_i, x_{test})$$

where y_{test} is the predicted value for a test example x_{test} . Using the same logic as in the kernel matrix reformulation, we get

$$y_{test} = \mathbf{K}_{eval} \alpha^*$$

where

$$\mathbf{K}_{eval} = \text{np.exp} \left(-\frac{1}{2\sigma^2} (\mathbf{B}^T - 2\mathbf{X}_{test}\mathbf{X}^T + \mathbf{C}) \right)$$

where \mathbf{X}_{test} is an $(m \times n)$ testing data matrix, \mathbf{B}^T is same as defined in 7 and \mathbf{C} is defined as the diagonals of matrix $\mathbf{X}_{test}\mathbf{X}_{test}^T$. Note however that the \mathbf{K}_{eval} matrix is not symmetric matrix like the kernel matrix; it is of dimensions $(m \times l)$.

2.2 Functions for regression experiments on Boston Housing Dataset

```

1 def naive_regression_eval(num_train_samples, num_test_samples, train_y,
2   test_y, results_dict):
3     """
4     Performs Naive Regression on given data and saves results.
5
6     Args
7     ----
8     num_train_samples : number of training samples. This will be used to
9       create a ones vector for the training set.
10    num_test_samples : number of testing samples. This will be used to
11      create a ones vector for the testing set.
12    train_y : NumPy array of training output values. Should be of shape (
13      num_train_samples, 1)
14    test_y : NumPy array of testing output values. Should be of shape (
15      num_test_samples, 1)
16    results_dict : Dictionary in which to save the results.
17
18    Returns

```

```

14  -----
15  None.
16
17  """
18
19  #Add the key into the results dictionary
20  key = "Naive Regression"
21  if not key in results_dict.keys():
22      results_dict[key] = {"Train MSE": [], "Test MSE": []}
23
24  #Create a ones vector for train and test data
25  train_ones = np.ones((num_train_samples, 1))
26  test_ones = np.ones((num_test_samples, 1))
27
28  #Compute the weights using training set - in case we encounter
29  singular matrices, Ive added the psuedoinverse as backup
30  try:
31      w = np.linalg.inv(train_ones.T @ train_ones) @ train_ones.T @
32      train_y
33  except np.linalg.LinAlgError:
34      w = np.linalg.pinv(train_ones.T @ train_ones) @ train_ones.T @
35      train_y
36
37  #Add results into the dictionary
38  results_dict[key]["Train MSE"].append(MSE(train_ones, train_y, w).
39  item())
40  results_dict[key]["Test MSE"].append(MSE(test_ones, test_y, w).item()
41  )
42
43  print(">> Naive Regression done.")
44
45  def linear_regression_eval(train_x, test_x, train_y, test_y,
46  results_dict, single=False):
47  """
48  Performs linear regression on the given data and saves results.
49
50  Args
51  ----
52  train_x : NumPy array of training inputs. Should be of shape (m, n)
53  where m is number of training examples and n is number of input
54  features
55  test_x : NumPy array of test inputs. Should be of shape (l, n) where
56  l is number of testing examples and n is number of input features
57  train_y : NumPy array of training outputs. Should be of shape (m, 1)
58  where m is number of training examples
59  test_y : NumPy array of testing inputs. Should be of shape (l, 1)
60  where l is number of training examples
61  results_dict : Dictionary in which to save the results.
62  single : Set to True if you wish to perform on single attributes.
63
64  Returns
65  -----
66  None.
67  """

```

```

58 if single:
59     num_features = train_x.shape[1] - 1
60
61     #Iterate over the features
62     for i in range(num_features):
63         #Create a key for the attribute and add it into the results
dictionary if it doesn't exist
64         key = "attribute {}".format(i+1)
65         if not key in results_dict.keys():
66             results_dict[key] = {"Train MSE": [], "Test MSE": []}
67
68         #Pick out a the specific feature and the bias column
69         train_x_subset = train_x[:, [0, i+1]]
70         test_x_subset = test_x[:, [0, i+1]]
71
72         #Checkpoint to make sure the dimensions are correct
73         assert train_x_subset.shape == (train_x.shape[0], 2)
74
75         #Computed least squares solution - In case we encounter singular
matrices, I have used the pseudoinverse as backup
76         try:
77             w = np.linalg.inv(train_x_subset.T @ train_x_subset) @
train_x_subset.T @ train_y
78         except np.linalg.LinAlgError:
79             w = np.linalg.pinv(train_x_subset.T @ train_x_subset) @
train_x_subset.T @ train_y
80
81         # Compute the train and test MSE and add it into the results_dict
82         results_dict[key]["Train MSE"].append(MSE(train_x_subset, train_y
, w).item())
83         results_dict[key]["Test MSE"].append(MSE(test_x_subset, test_y, w
).item())
84
85         print(">> Linear regression for {} done.".format(key))
86
87 else:
88     #Create a key for the attribute and add it into the results
dictionary if it doesn't exist
89     key = "all attributes"
90     if not key in results_dict.keys():
91         results_dict[key] = {"Train MSE": [], "Test MSE": []}
92
93     #Compute the weights using training set - in case we encounter
singular matrices, Ive added the psuedoinverse as backup
94     try:
95         w = np.linalg.inv(train_x.T @ train_x) @ train_x.T @ train_y
96     except np.linalg.LinAlgError:
97         w = np.linalg.pinv(train_x.T @ train_x) @ train_x.T @ train_y
98
99     #Compute the train and test MSE and add it into the results
dictionary
100     results_dict[key]["Train MSE"].append(MSE(train_x, train_y, w).item
())
101     results_dict[key]["Test MSE"].append(MSE(test_x, test_y, w).item())
102

```

```

103     print(">> Linear Regression for {} done.".format(key))
104
105 def ridge_regression_eval(train_x, train_y, test_x, test_y,
106     sigma_values, gamma_values, results_dict):
107     """
108     Performs kernelized ridge regression on the given data and saves
109     results.
110
111     Args
112     ----
113     train_x : NumPy array of training inputs. Should be of shape (m, n)
114               where m is number of training examples and n is number of input
115               features
116     test_x : NumPy array of test inputs. Should be of shape (1, n) where
117               1 is number of testing examples and n is number of input features
118     train_y : NumPy array of training outputs. Should be of shape (m, 1)
119               where m is number of training examples
120     test_y : NumPy array of testing inputs. Should be of shape (1, 1)
121               where 1 is number of training examples
122     sigma_values : NumPy 1D array of possible sigma values.
123     gamma_values : NumPy 1D array of possible gamma values.
124
125     Returns
126     ----
127     None.
128     """
129
130     #Generate a seed value which we will use inside the for loop.
131     seed=np.random.randint(100)
132
133     #Create a key for the attribute and add it into the results
134     dictionary if it doesn't exist
135     key = "Kernel Ridge Regression"
136     if not key in results_dict.keys():
137         results_dict[key] = {"Train MSE": [], "Test MSE": []}
138
139     #The best parameters will be saved in the order [best_sigma,
140     best_gamma]
141     best_parameters = {"sigma":0, "gamma":0}
142
143     #Initialize the smallest error to be some large number. This will be
144     used to update our best parameters
145     min_error = 10000000000000000000
146
147     #Iterate over entire parameter space
148     for i in range(len(sigma_values)):
149         for j in range(len(gamma_values)):
150
151             #Create the sigma and gamma pair
152             sigma = sigma_values[i]
153             gamma = gamma_values[j]
154
155             #Perform 5-fold cross validation.
156             # Note - The train_with_kfoldCV() function performs a kfoldCV for
157             just one sigma and gamma.

```

```

147     # In the function, I've added a shuffle operation BEFORE
    splitting the data into the k groups.
148     # This means that everytime we use a new sigma and gamma, the
    group compositions may differ due to the shuffle.
149     # We need to ensure that the group compositions after shuffling
    are the same everytime so that our parameters sigma and gamma are
    comparable.
150     np.random.seed(seed)
151     mean_cv_error = train_with_kfoldCV(k=5, data=(train_x, train_y),
    kernel_params=(sigma, gamma), shuffle=True, verbose=False)
152
153     #If the mean_cv_error is smaller than the minimum error, then we
    can update our best parameters
154     if mean_cv_error < min_error:
155         min_error = mean_cv_error
156         best_parameters['sigma'] = sigma
157         best_parameters['gamma'] = gamma
158
159     #Compute dual regression solution using training set
160     alpha_star = kernel_ridge_regression(X=train_x, y=train_y, sigma=
    best_parameters['sigma'], gamma=best_parameters['gamma'])
161
162     #Compute predictions on train set: The "test" points here are just
    the training set again
163     train_y_pred = evaluate(alpha=alpha_star, X_tr=train_x, X_te=train_x,
    sigma=best_parameters['sigma'])
164     #Compute predictions on test set
165     test_y_pred = evaluate(alpha=alpha_star, X_tr=train_x, X_te=test_x,
    sigma=best_parameters['sigma'])
166
167     #Compute train and test MSE
168     train_MSE = (1/len(train_y)) * (train_y_pred - train_y).T @ (
    train_y_pred - train_y)
169     test_MSE = (1/len(test_y)) * (test_y_pred - test_y).T @ (test_y_pred
    - test_y)
170
171     results_dict[key]["Train MSE"].append(train_MSE)
172     results_dict[key]["Test MSE"].append(test_MSE)
173
174     print(">> Kernel Ridge Regression done.\n")
175
176     #Initialize a final results dictionary which we shall update with all
    the errors
177     final_results = dict()
178
179     #Store number of runs here
180     runs = 20
181
182     #Generate a seed for reproducible results
183     np.random.seed(1753295)
184
185     #PROBLEM: INSPITE OF SETTING SEED HERE, THE FUNCTION IS GENERATING SAME
    TRAINING SET EVERYTIME!! HOW TO FIX IT?
186
187     #Iterate over runs

```

```

188 for i in range(runs):
189
190     print("Run {}".format(i+1))
191
192     #Generate a fresh batch of train/test data
193     train_X, train_Y, test_X, test_Y = split_data(inputs=X, targets=Y,
194         test_proportion=1/3, shuffle=True)
195     # print("Training set = ", train_X)
196
197     #PERFORM NAIVE REGRESSION AND UPDATE RESULTS
198     naive_regression_eval(num_train_samples=len(train_X),
199         num_test_samples=len(test_X), train_y=train_Y, test_y=test_Y,
200         results_dict=final_results)
201
202     #PERFORM LINEAR REGRESSION ON SINGLE ATTRIBUTES AND UPDATE RESULTS
203     linear_regression_eval(train_x=train_X, test_x=test_X, train_y=
204         train_Y, test_y=test_Y, results_dict=final_results, single=True)
205
206     #PERFORM LINEAR REGRESSION ON ALL ATTRIBUTES AND UPDATE RESULTS
207     linear_regression_eval(train_x=train_X, test_x=test_X, train_y=
208         train_Y, test_y=test_Y, results_dict=final_results, single=False)
209
210     #PERFORM KERNELIZED RIDGE REGRESSION ON ALL ATTRIBUTES AND UPDATE
211     RESULTS
212     #Create parameter space
213     sigma_values = 2*np.arange(7, 13.5, 0.5)
214     gamma_values = (2*np.arange(40.0, 25.0, -1))**-1
215     # sigma_values = 2*np.arange(7, 8.0, 0.5)
216     # gamma_values = (2*np.arange(40.0, 38.0, -1))**-1
217     ridge_regression_eval(train_x=train_X[:, 1:], train_y=train_Y,
218         test_x=test_X[:, 1:], test_y=test_Y,
219         sigma_values=sigma_values,
220         gamma_values=gamma_values,
221         results_dict=final_results)
222
223     print("Evaluation Completed.")

```